

A correction testbed with distinct controls

Three realizations × five dependent orders; endpoint depends on dataset

DRAFT — lead review

Learned-reader condition

Frozen GPT-2 base

Fixed BP-trained reader

Adaptive bounded memory

Controlled comparisons

Random reader / stable cap

Live cap / matched update

S1 continued-base controls

Separate questions

Acquisition and retention

Locality and near misses

Ordinary-text probability harm

zsRE / CounterFact: 1,000 edits. MQuAKE: descriptive at 300. This reader study uses BP, not PC credit.

50 locality prompts do not exhaust specificity. Five orders within a realization are dependent.

Claims: R1-design, R1-controls, R1-fidelity