


Predictive coding: retention, harm and learning cost

Depth controls: order100, three realizations per dataset; exposed historical streams.

DRAFT — lead review

More settling: better taught-answer retention, higher cost and ES99+
zsRE · same order 100 · three realization means



Exposed S5. 1,000 edits; harm on 4,064 fixed positions. One-step endpoint scores equal adjoint; harm differs slightly numerically.
No confidence interval or depth selection. CounterFact is saturated / at a paraphrase floor; full results in the report.

Random: 3/993 immediate successes. Adjoint budget offered, underspent.

More own-prompt retention does not establish better generalization or equal-compute superiority.

Claims: PC-depth, PC-one-step, PC-random, PC-budget