

What we learned, what remains, what it took

Measured credit and tail interventions; autonomous audit policies remain proposed.

DRAFT — lead review

Active inference

Concrete correction testbed

Autonomous audit policy

REMAINS PROPOSED

Predictive coding

Credit: harm/cost trade-offs

Reader: 3 seeds; $\sim 100\times$ cost

CF PARAPHRASE LOSS: ALL SEEDS

Extremes

Mixture bounds token harm

All ten memories pass rule

MEASURED; FAMILY UNPROVEN

270-cell snapshot: 392.42 process-hours, not elapsed GPU hours. PC costs reported separately.

GPT-2 small (124M). GPU work complete Oct 4; review freeze Oct 9; talk Oct 15. Upper-layer results in backup.

Claims: AI-next, PC-depth, PC-budget, AW-B, HT13-corrected, talk-scope, model-scale, PC-reader, R-extension, AW-L