

11 · Learning the reader improves paraphrase retention

Retained paraphrase score, averaged over orders within each realization, then over the three realizations. Higher is better.

Dataset and endpoint	Learned reader	Random reader	Stable v0 cap
zsRE · 1,000 edits	96.0%	52.4%	18.6%
CounterFact · 1,000 edits	67.8%	12.1%	0.0%
MQuAKE · 300 edits	71.6%	0.0%	0.0%

The learned-reader realization means span 95.5–96.5% on zsRE, 66.6–69.2% on CounterFact and 70.3–72.3% on MQuAKE.

The cap often carries an answer beyond its teaching prompt. The strength of that transfer depends on the dataset.

Source: Stage-4 triplet summary. Bars start at zero on a common 0–100% scale; ranges are descriptive.