

16 · The mixture reduced severity to about one-third

Mean loss increase among positions with $\Delta > 0.01$ nat, on ten exposed memories: two datasets × five orders at 300 edits.

Dataset	Original cap	Probability mixture
zsRE	1.619 nats	0.542 nats
CounterFact	1.925 nats	0.581 nats

All ten memories met the declared retention-and-preservation rule. Paraphrase retention was unchanged on zsRE and fell by 0.67–1.17 percentage points on CounterFact.

Harm frequency barely changed. CounterFact mean KL still exceeded 0.001.

This is a measured reduction in the severity of unintended changes, with a small editing trade-off.